

STITCH FIX

# Spark SQL Performance Troubleshooting and Tuning

Sky Yin

Jun 2, 2016

# Quick Start

---

## Spark is a **distributed computing platform**

- Distributed  $\neq$  Scalable  $\neq$  Easy to scale
  - Redshift vs Spark
- Not for data storage (except caching)
- 4 horse chariot
  - 4 language support: Java, Scala, Python, R
  - 4 use cases:
    - **SQL**
    - Machine learning
    - Graph analysis
    - Stream processing

# Quick Start

---

## Two ways to run a Spark SQL query

- Flotilla
  - PySpark notebook (not ideal, better to use Presto)
- ETL command
  - Demo 1: [https://github.com/stitchfix/quant\\_workshop/tree/master/sparksql](https://github.com/stitchfix/quant_workshop/tree/master/sparksql)
  - How does it run? ETL => Genie => Spark cluster

# Quick Start

---

## Resources

- Our [Redshift -> Spark migration guide](#) (thanks Doug, Daragh, Matt, Sky and etc.)
- Spark SQL function reference
  - [pyspark.sql.functions module](#) section in PySpark doc
- [Stack Overflow #apache-spark-sql](#)
- Slack channel #spark-users

# Understanding Spark SQL Execution

---


## Fundamental building block: RDD

- Resilient Distributed Dataset
- RDD partition  $\neq$  Hive metastore partition
- Everything is RDD
  - SQL query
    - => a graph of RDD transformation and action
    - => Output

# Understanding Spark SQL Execution

**Key question: logical plan => physical plan?**

- Demo 2
- Spark execution model
- Physical plan
  - Jobs => Stages => Tasks



# Monitoring Spark SQL Performance

## Spark console

- [go/spark-console](#) (thanks Tarek)




### All Applications

#### Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooter Nodes
10284	0	9	10275	83	248.38 GB	686.25 GB	0 B	83	247	0	30	0	0	181	0	
Show 20 entries Search:																
ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI						
<a href="#">application_1456900025701_8690</a>	trached	pyspark-shell	SPARK	default	Fri, 27 May 2016 22:33:43 GMT	N/A	RUNNING	UNDEFINED	<div></div>	<a href="#">ApplicationMaster</a>						
<a href="#">application_1456900025701_8220</a>	uraza	pyspark-shell	SPARK	default	Thu, 26 May 2016 21:56:05 GMT	N/A	RUNNING	UNDEFINED	<div></div>	<a href="#">ApplicationMaster</a>						
<a href="#">application_1456900025701_6293</a>	jmagнусson	pyspark-shell	SPARK	default	Tue, 17 May 2016 17:12:39 GMT	N/A	RUNNING	UNDEFINED	<div></div>	<a href="#">ApplicationMaster</a>						
<a href="#">application_1456900025701_5693</a>	jacobperkins	pyspark-shell	SPARK	default	Mon, 09 May 2016 19:35:57	N/A	RUNNING	UNDEFINED	<div></div>	<a href="#">ApplicationMaster</a>						

# Monitoring Spark SQL Performance

## Spark tracking UI



JobsStagesStorageEnvironmentExecutorsSQL

quant-hour-sparksql application UI

### Spark Jobs (?)

Total Uptime: 2.2 min  
Scheduling Mode: FIFO  
Active Jobs: 1  
[Event Timeline](#)

#### Active Jobs (1)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	<a href="#">foreachPartition at /mnt/tomcat/genie-jobs/7811ce94-edea-4a59-92d2-c624671770fb/demo2.py:14</a>	2016/06/02 05:30:10	1.2 min	0/2	<div>978/1201</div>



# Monitoring Spark SQL Performance

[Jobs](#)[Stages](#)[Storage](#)[Environment](#)[Executors](#)[SQL](#)

quant-hour-sparksql application UI

## Details for Job 0

**Status:** RUNNING**Active Stages:** 2**Pending Stages:** 2

- ▶ [Event Timeline](#)
- ▶ [DAG Visualization](#)



### Active Stages (2)

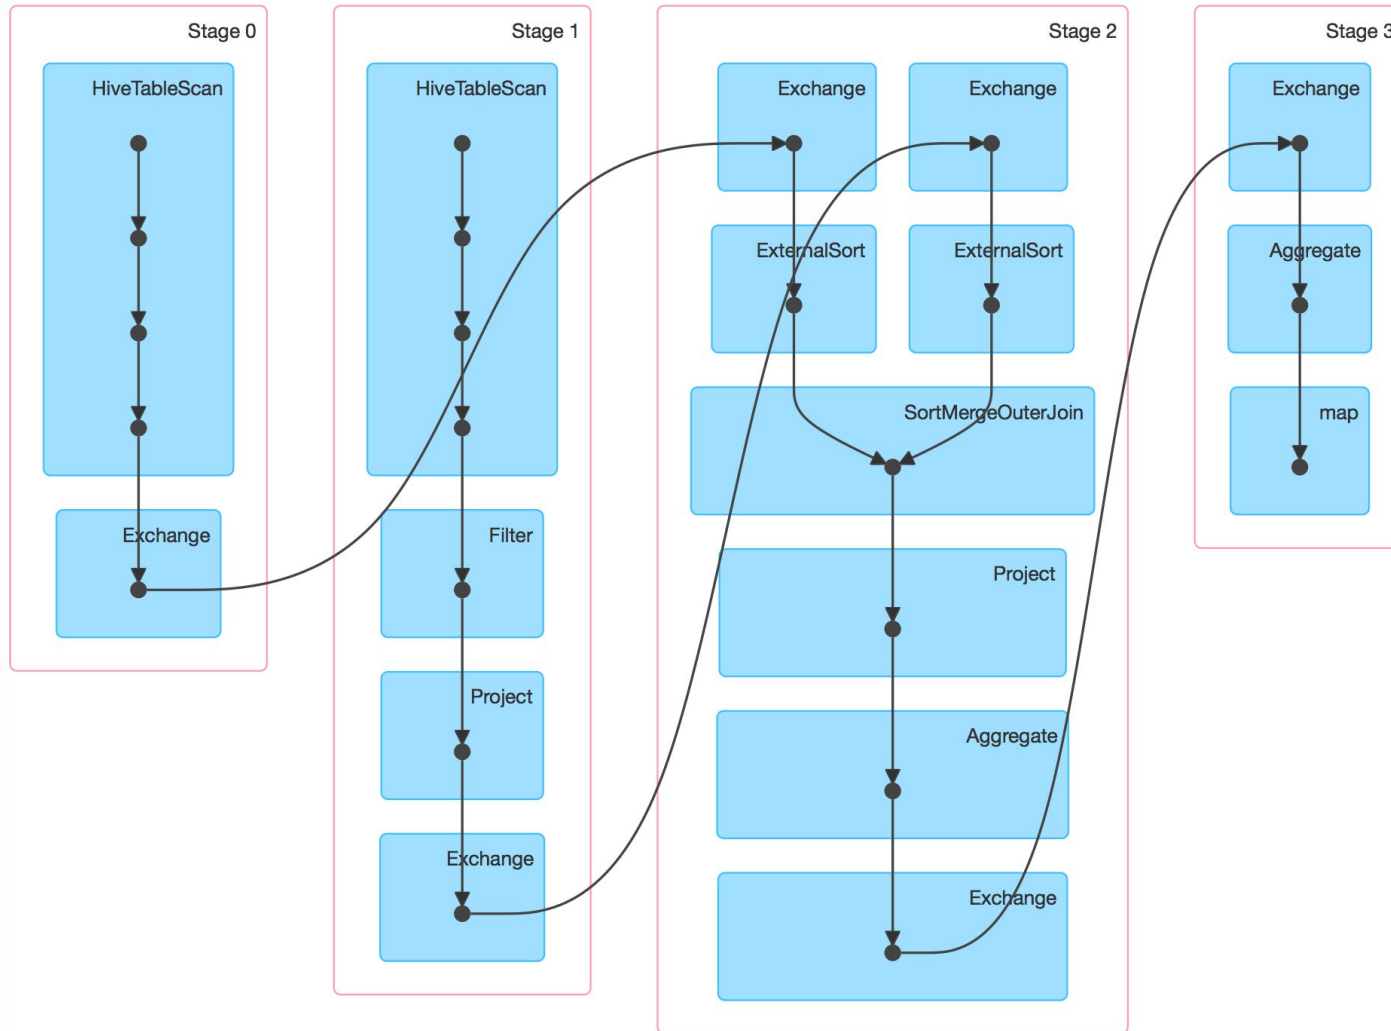
Stage Id	Description			Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	<a href="#">showString at NativeMethodAccessorImpl.java:-2</a>	<a href="#">+details</a>	(kill)	2016/06/02 08:39:37	1.1 min	0/2				
0	<a href="#">showString at NativeMethodAccessorImpl.java:-2</a>	<a href="#">+details</a>	(kill)	2016/06/02 08:39:37	1.1 min	150/200	108.4 MB			116.8 MB

### Pending Stages (2)

Stage Id	Description			Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
3	<a href="#">showString at NativeMethodAccessorImpl.java:-2</a>	<a href="#">+details</a>		Unknown	Unknown	0/1				
2	<a href="#">showString at NativeMethodAccessorImpl.java:-2</a>	<a href="#">+details</a>		Unknown	Unknown	0/200				

# Monitoring Spark SQL Performance

## ▼ DAG Visualization



# Monitoring Spark SQL Performance

## Details for Stage 0 (Attempt 0)

**Total Time Across All Tasks:** 4.7 min

**Input Size / Records:** 145.6 MB / 11348443

**Shuffle Write:** 155.8 MB / 11348443

- ▶ [DAG Visualization](#)
- ▶ [Show Additional Metrics](#)
- ▶ [Event Timeline](#)

## Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration		0.6 s	0.7 s	0.9 s	42 s
Scheduler Delay	4 ms	7 ms	9 ms	12 ms	0.1 s
Task Deserialization Time	2 ms	4 ms	5 ms	10 ms	4 s
GC Time		0 ms	0 ms	0 ms	0.6 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms	0 ms	0 ms	0 ms	0 ms
Peak Execution Memory	0.0 B	0.0 B	0.0 B	0.0 B	0.0 B
Input Size / Records	0.0 B / 56022	758.5 KB / 56580	760.4 KB / 56735	762.5 KB / 56891	769.7 KB / 57550
Shuffle Write Size / Records	788.9 KB / 56022	795.3 KB / 56580	797.6 KB / 56735	799.8 KB / 56891	809.5 KB / 57550

# Tuning Spark SQL Performance

---

## Parallelism

- Number of executor `--spark_num_executors`
- Number of cores/executor `--spark_executor_cores`
- Number of partition
  - Spark will run one task for each RDD **partition**
  - In PySpark `rdd.getNumPartitions()`
- Repartition

# Tuning Spark SQL Performance

---

## Memory

- driver vs worker
  - don't call `collect()` on a large query result
- executor memory
  - executor loss
  - `java.lang.OutOfMemoryError: GC overhead limit exceeded`
- GC
- Caching
  - space/speed tradeoff

# Tuning Spark SQL Performance

---

## Join performance

- Try to avoid shuffle
  - ShuffledHashJoin vs BroadcastHashJoin
  - `spark.sql.autoBroadcastingJoinThreshold`
- Filter before join
- Demo 3
  - [SortMergeJoin](#)
- Detect uneven shuffle
  - Tasks taking much longer to run
  - Tasks with much higher input or shuffle output
- Think about the actual distribution of data
  - Ex. group by gender?